

Distributed Cooperative Robot Systems

Students of the Lecture

December 18, 2016

Contents

1	Introduction	3
1.1	Purpose of this Document	3
1.2	How to Contribute to this Document?	3
2	External Software and Developing Tools	5
2.1	GIT and GitHub	5
2.2	Robot Operating System (ROS)	5
2.3	Build Chain	6
3	Process Manager and Remote Control GUI	7
3.1	Quickstart Guide	7
3.2	General Information on the Process Manager	8
3.3	General Information on the Remote Control GUI	9
3.4	Configuration	11
3.5	Future Work	13
4	ALICA Client	14
5	TurtleBot Hardware	15
6	TurtleBot Software	16
6.1	Quick Start Guide	16
6.2	Drive Topic and Message	18
7	TurtleBot Gazebo Simulation	19
7.1	Starting Gazebo simulation	19
7.2	.bashrc	19

1 Introduction

1.1 Purpose of this Document

Instead of giving the students a lot of Powerpoint slides that compress the complex topics of this course into a bunch of bullet points, we wanted the students to have a dedicated document for preparing their exam. Unfortunately, such a document is a lot more work and it needs several iterations of thorough reviewing. Therefore, we came up with the compromise that the students write this document themselves and the lecturers only have to review and improve it.

Maybe you ask: "What is wrong with Powerpoint?"

A Powerpoint presentation should aid a presenter in presenting a certain message or information to his audience. The slides alone often miss critical parts (that couldn't easily condensed into bullet points) and are unclear without the corresponding vocal line of the presenter. Nevertheless, students end up alone with the slides trying to prepare their examinations.

Apart from that, try to google topics like "disadvantages of powerpoint" or "why we should ban Powerpoint" and you will find interesting articles that connect Powerpoint and the Accident of the Columbia Space Shuttle [1] [2].

1.2 How to Contribute to this Document?

Basically there are two aspects. At first you need to get the L^AT_EX source code of this document from our cnc-turtlebots GitHub Repository. It is located in the folder doc/latex. Anybody can download it, but for making changes to it you need to have an GitHub-Account which was given the corresponding privileges. That is typically arranged during the first session of this course.

This leaves us with the requirements, that you are able to write L^AT_EX and know how to use GIT in combination with GitHub. Regarding GIT and GitHub we recommend to read Section 2.1.

Maybe L^AT_EX is hard to master when it comes to tables, graphics and customizing the layout of a document, but you don't have to do that for contributing to this document. You should only be able to write text referencing images and place useful links from time to time. Furthermore, you should be able to structure the content into chapters, sections, and subsections. Everything just mentioned is already done in this document, so you can start by copying commands from the corresponding sections.

For further reference and an in-depth study of L^AT_EX, we recommend the following sources:

L^AT_EX – eine Einführung und ein bisschen mehr: A little bit longer introduction to L^AT_EX basics.

L^AT_EX WikiBook: Nice for looking up simple stuff.

tex.stackexchange.com: Stack Overflow for L^AT_EX

TikZ: L^AT_EXpackage for designing advanced graphics directly in L^AT_EX.

The final issue is to decide which editor you choose for compiling this L^AT_EXdocument. We recommend TexMaker, because it is available for Ubuntu and Windows, although, there currently exists an issue with short cuts under Ubuntu 16.04 LTS. Consider this Ask Ubuntu Post for solving it.

With TexMaker you only have to open the `Software_Reference_Book.tex` file and click on **Quick Build**. If you get compile errors, you probably need to install the necessary L^AT_EXpackages. Under Ubuntu this is done by executing the following console command:

```
sudo apt-get install texlive-full
```

2 External Software and Developing Tools

This chapter is intended to give a short introduction to external software packages or frameworks, we utilise for the Carpe Noctem Cassel software framework.

2.1 GIT and GitHub

GIT is one of the most advanced version control systems currently available. Nevertheless, during our daily work we only use 20% of its functionality. So for starters try to learn the stuff you need and ignore its advanced features like *rebase* and *cherry pick*.

The best reference and documentation about GIT can be found at <http://www.git-scm.com/docs>

Most of our software is published open source under the MIT License at our GitHub Repositories. Therefore, it is also interesting to read about the features of GitHub, like SSH-Key based authorization, groups, organisations, and MarkDown.

The README.md files in our repositories are written in MarkDown, because GitHub parses these MarkDown files and auto-generates an HTML documentation from it.

A list of GIT commands, that should be enough for the start, can be found on git-tower.com.

2.2 Robot Operating System (ROS)

ROS, as we use it, is a simple inter process communication middleware. Before you ask, yes it is not intended to be used for inter machine/robot communication. Therefore, we have developed a simple ROSUdpProxy, for our purposes at RoboCup.

Tutorials for ROS can be found here: <http://wiki.ros.org/ROS/Tutorials>

After following this tutorial you should be able to explain a bunch of things:

Topics: How do they work?

Nodes: What is a ROS node?

roscore: What is its job?

package.xml: build_depend, run_depend, licences, ...

CMakeLists.txt: What are the critical ROS specific macros and how do they work?

Console Commands: rosrun, rospack, rosdep, rosdep, rosdep, roslaunch, rostopic, rosnode

catkin_make: How to compile a workspace/a package?

ROS-Workspace: What is its structure and why is it structured that way?

ROS-Services: How are they defined, compiled/generated, and how do they work?

ROS-Messages: How are they defined and compiled/generated?

roscpp API: How to create a publisher and a subscriber?

2.3 Build Chain

We utilise *catkin* from the ROS Universe as our build chain. Catkin is basically a workspace-oriented extension of CMake. Therefore, it heavily relies on CMake and in order to understand catkin it is recommended to understand CMake first.

CMake is open source and developed by KitWare. Basically CMake autogenerates Makefiles out of CMakeLists.txt files located in each software module. Therefore, our build chain can really be considered as a chain ⊕:

Catkin $\xrightarrow{\text{manages}}$ CMake $\xrightarrow{\text{auto-generates}}$ Makefiles $\xrightarrow{\text{commands}}$ GCC $\xrightarrow{\text{to compile}}$ executables and libraries.

3 Process Manager and Remote Control GUI

The process manager is an executable for managing the processes running on a PC. Compared with the former C# framework, it replaces the process manager Care. The process manager can run on the robot for testing with real a robot, or on your local PC for testing with a simulator (add `-sim` as parameter). With the help of the *ROBOT* environment variable, it is possible to manage processes for multiple robots on a single PC, which is useful for multi-robot testing on a single PC. The name of the ROS-Package of the process manager is *process_manager* and can be found by using `roscd process_manager`.

The remote control GUI for the process manager is an RQT plugin (see <http://wiki.ros.org/rqt> for details). In the former C# framework this GUI was highly integrated into the LebtClient and mixed with robocup msl specific GUI elements. The idea of the new GUI is, to make it useable in other domains, too. The name of the ROS-Package of the remote control GUI is *pm_control* and can be found by using `roscd pm_control`.

3.1 Quickstart Guide

In order to bring up a single process manager and its remote control GUI, on the same PC, execute the following commands in the given order:

- `rosrun process_manager process_manager`
- `rosrun pm_control pm_control`

The first command starts the process manager, which will automatically start a roscore, if none is running. Please add the `-sim` parameter to the process manager, if you want to use it locally managing multiple robots (e.g. for simulation). The second command start the remote control GUI, which should display the received information of the process manager (see Figure 3.1).

If the process manager and the remote control GUI should run on different PCs, which are in the same network, you need to make sure, that on both machines a roscore and a UDP proxy is running. Therefore, you need to execute the following commands on the machine, where the process manager should run:

- `rosrun msl_udp_proxy msl_udp_proxy`
- `rosrun process_manager process_manager`

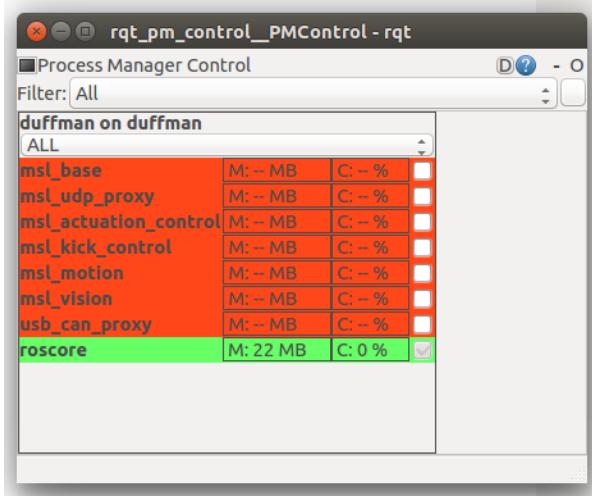


Figure 3.1: The Remote Controle GUI

The msl-udp-proxy is an autogenerated proxy, which broadcasts MSL specific ROS messages from the local machine to a multicast address on the network. You can also use another UDP proxy (e.g. ttb-udp-proxy, or bbb-udp-proxy), as along as it forwards the *ProcessCommand* and *ProcessStats* ROS messages of the process manager.

On the machine, where the remote control GUI should run, execute:

- `roscore`
- `rosrun msl_udp_proxy msl_udp_proxy`
- `rosrun pm_control pm_control`

You need to start the roscore manually, as the remote control GUI does not start a roscore automatically.

3.2 General Information on the Process Manager

PROC Filesystem

The information about the managed processes are collected from the */proc* filesystem. This is a specific path, which is available on each linux-kernel-based operating system. Inside the */proc* folder, each running process has a seperate subfolder, named after its process id (PID). Please note, that this PID is determined by the kernel and has nothing to do with the process ids of the *ProcessManaging.conf* file. The process manager uses the information available in the processes *cmdline* and *stat* file (e.g. */proc/2341/cmdline*). These files are continuously updated by the kernel. For

further information about the proc filesystem, consider chapter 7 of the book *Advanced Linux Programming* (available at: <http://www.advancedlinuxprogramming.com/alp-folder/>).

The interesting thing about the proc filesystem approach is, that it enables the process manager to attach to processes it did not start in the first place! If you had some processes already up and running, just start the process manager and its remote control GUI to see the statistics about those processes.

Another fact about the process manager is, that it does not do anything to the processes, when it is closed. The launched processes are independent of the process manager and continue their execution without it. You accidentally stopped the process manager? No problem just restart the process manager and it will continue to monitor the started processes.

Communication

The process manager defines two ROS messages: *ProcessCommand* and *ProcessStats*. The *ProcessCommand* message is used to let him start and stop processes. The *ProcessStats* message is used by the process manager to report the statistics about its managed processes. The process manager is configured to scan the proc filesystem every two seconds, so that the *ProcessStats* is send roughly every two seconds.

The process manager subscribes on the ROS topic `/process_manager/ProcessCommand` and publishes its commands to `/process_manager/ProcessStats`.

Allowed Number of Processes

The process manager is written in a way, that it only allows to run a certain process only one time per robot. So if you want to run a process twice on the same machine, you need to modify the *ROBOT* environment variable for at least one of the two process instances. E.g.: `ROBOT=nase rosrun msl_base msl_base -m WM16` and `ROBOT=myo rosrun msl_base msl_base -m WM16`.

Another feature of the limitation on the number of allowed processes is, that if you did start too many processes, e.g. two or more image processing processes on one robot, you can simply start the process manager and it will clean up the mess. It will kill all but one process of each kind and start reporting statistics about the left processes.

3.3 General Information on the Remote Control GUI

Process GUI

The remote control GUI has a single process GUI element for each process (see Figure 3.2). From left to right it shows: the process name, its memory usage in MB, its cpu usage in percent (100% means one core), its check box for start and stop. The background color of the process GUI is either red (not running), green (running), gray (unknown). Note that the check box does not determine, whether a process is

running or not, it is just for starting and stopping a process. Start-Commands for running processes are ignored by the process manager. The check box is disabled for the roscore, because stopping it would cut the communication to the process manager. Check boxes of other processes are disabled too, if they are running with parameters that differ from the currently selected bundle (see 3.3).



Figure 3.2: GUI Elements for one Process

If you have a running process (background is green), you can hover over the process GUI element, in order to show its ToolTip. The ToolTip shows the command, which was used to start the process. This way, you don't have to check the ProcessManaging.conf file, in order to know what parameters are used in a certain bundle.

Communication

The GUI elements are created at runtime, when a corresponding message from a process manager arrives. So if you don't receive any ProcessStats messages, your remote control GUI won't show anything. Furthermore, the GUI elements are deleted, if you don't receive messages for certain amount of time (roughly 3 seconds). The remote control GUI subscribes on the ROS topic `/process_manager/ProcessStats` and publishes its commands to `/process_manager/ProcessCommand`.

Bundle Selection

Selecting a bundle in the drop down box of the GUI, means that the listed processes will be started with the corresponding parameter set, as specified in the ProcessManaging.conf file (see Section 3.4). It also means, that you cannot stop a process which runs with a different parameter set, than specified in the selected bundle. In such cases, the check box of the process is disabled (grayed out). Nevertheless, there are two default bundles: ALL and RUNNING. If you select one of these two bundles, you can interact with all processes.

ALL lists all processes configured in the ProcessManaging.conf file. This is useful, if you want to start a set of processes, which is not specified as an explicit bundle. RUNNING lists all processes of the ProcessManaging.conf file, which are currently up and running. This is useful, for determining, whether there are unwanted processes running, which are not part of the bundle that you would like to use.

The `roscore` process is specially handled. If the roscore is stopped, the process manager cannot receive commands anymore. Therefore, it is not allowed to stop the roscore process within the remote control GUI.



Figure 3.3: Bundle Selection by Drop Down Box

3.4 Configuration

In order to configure, which processes should be managed by the process manager, you need to edit the *ProcessManaging.conf* file. The file is usually located in the *etc* folder, determined by the *DOMAIN_CONFIG_FOLDER* environment variable. Inside the *ProcessManaging.conf* file several comments explain the config values itself. Nevertheless, lets explain the config for the *msl_base* process in detail:

```
[Base]
id = 7
execName = msl_base
rosPackage = msl_base
mode = none
[paramSets]
1 = -m, TwoHoledWallMaster
2 = -m, ActuatorTestMaster
3 = -m, WM16
4 = -m, TestApproachBallMaster
5 = -m, TestCheckGoalKicki
6 = -m, WM16, -sim
[!paramSets]
[!Base]
```

The **name** (Base) of the outmost section denotes the GUI-String representing the executable. The executable denoted by **execName** (*msl_base*) need to be located in the *PATH* environment variable, or should be found by executing
`catkin_find --libexec msl_base`

in order to work with the process manager. Here *msl_base* is the **rosPackage** name.

The **id** of the process must be unique in the *ProcessManaging.conf* file and is used to refer to this process in the bundles section (explained later) and in the messages send to and received from the process manager.

The **mode** decides how the process manager handles crashes of the process and some other things. At the moment, there are 3 different modes.

none Basically does nothing. It does not autostart the process, when the process manager is started with the *-autostart* parameter. It does not restart the process, when it did crash.

keepAlive Processes, configured with this mode, will be restarted by the process manager, when they crashed.

autostart This mode makes the process manager start this process, when the process manager is started with the *-autostart* parameter and restarts it after crashes.

In the **paramSets** sections, it is possible to specify different sets of parameters which can be used to start the process. Each parameter set follow a simple key-value-pair convention, where the key must be a (for this process unique) positive integer, greater than 0. The value is a comma (,) separated list of parameters. Please note, that the parameter **-m TwoHoledWallMaster** is actually two parameters: **-m** and **TwoHoledWallMaster**. The parameter set with the lowest key is considered to be the default parameter set, which means that this parameter set is used, if not specified otherwise (e.g. by choosing a bundle).

The **bundles** section of the ProcessManaging.conf file allows to specify a set of processes with a specific parameter set for each of it. Here is a small example:

```
[Bundles]
[WM16]
    processList      = 0,1,2,3,4,5,6,7
    processParamsList = 0,0,1,0,0,0,0,3
[!WM16]

[TwoHoledWall]
    processList      = 0,1,2,3,4,5,6,7
    processParamsList = 0,0,1,0,0,0,0,1
[!TwoHoledWall]
[!Bundles]
```

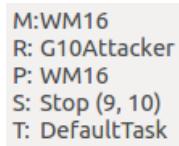
In this example, two bundles are specified: WM16 and TwoHoledWall. Each bundle consists of two key-value-pairs: **processList** and **processParamsList**. The **processList** specifies the list of processes, by listing the process ids. The **processParamsList** specifies the parameter sets for each process. It is important that the order of the **processParamsList** has to be the same as in the **processList**. For example: In the WM16 bundle the process with the id 7 (Base) is started with the parameter set id 3 (**-m, WM16**). In the TwoHoledWall bundle, it is started with the parameter set id 1 (**-m, TwoHoledWallMaster**). The parameter set id 0 is a special value, with the meaning, that there is no parameter set specified for this process, because it has no parameters. Another special value, which is only used in the messages send to and received from the process manager, is -1. It says that the parameters of a process are unknown, e.g., the **msl_base** is running with an unknown master plan (**-m FancyNewTestPlanMaster**).

3.5 Future Work

- Make the retry timeout for starting processes a parameter in the ProcessManaging.conf file.
- Make it possible to add the GUI for another virtual robot, in order to command a process manager to start processes for another robot. This is necessary for local testing with multiple robots.
- Document the implemented feature of starting launch scripts here. (For details see the ProcessManaging.conf file).

4 ALICA Client

The ALICA Client is a simple GUI for visualising AlicaEngineInfo messages. Currently it can be started for visualising the messages of a single engine by calling `rosrun alica_client alica_client`. If you have several robots running, the GUI will flicker between their incoming messages. For visualising the messages of multiple robots you should use the Robot Control GUI (see ??). It integrates the same GUI component multiple times.



M:WM16
R: G10Attacker
P: WM16
S: Stop (9, 10)
T: DefaultTask

Figure 4.1: The ALICA Client GUI

In Figure 4.1 you can see the ALICA Client GUI. **M:** denotes the currently executed master plan (WM16). **R:** denotes the current role of the robot (G10Attacker). **P:** is the deepest plan in the alica plan hierarchy, the robot is currently executing. As it is WM16, it means, that it is inside the master plan, but not deeper. Inside this plan, the robot is currently one of two robots inside the Stop state (**S:**). In this case, it is either the robot with id 9 or 10. The name of the task associated with the active state is the DefaultTask (**T:**).

This tool is definitely usefull for debugging ALICA plans, but please note, that the AlicaEngineInfo messages are only send roughly every 100ms, but the plan state of an ALICA Engine typically changes much faster, than that. Therefore, you want recognize agents racing through the plans' state machines. Making this visible needs a little bit more sophisticated GUI, which should be able to read logs of the ALICA engine itself. Please see the project description for some details, about this possible bachelor project.

5 TurtleBot Hardware

6 TurtleBot Software

The TurtleBot Software consists of several launch files which start the processes and load the needed xml descriptions.

6.1 Quick Start Guide

In order to use a laptop to control a TurtleBot several environment variables have to set in the .bashrc, which are show in the following List 6.1 for Robot Leonardo.

- `export ROBOT=leonardo`
- `export TURTLEBOT_3D_SENSOR=kinect`
- `export TURTLEBOT_MAP_FILE="$DOMAIN_CONFIG_FOLDER/map_final.yaml"`

The export \$DOMAIN_CONFIG_FOLDER should have been set already by the setup script (ttb-setup.sh) found in cnc-turtlebots/configuration/setupscripts. Otherwise it has to be set to cnc-turtlebots/etc. Once these preparations are done the laptop can be used to control it should be placed with closed lid on the layer below the laser scanner. After this is done the USB-hub can be connected with the laptop and the TurtleBot is ready. The setup of the Turtlebot is shown in Figure 6.1. To prepare the laptop/pc



Figure 6.1: TurtleBot setup.

to communicate with the TurtleBot these processes are started in three terminals are need on the local laptop/pc in which the following commands have to be executed:

- `roscore`

- `rosrun ttb_udp_proxy ttb_udp_proxy`
- `roslaunch turtlebot Bringup rviz_robot.launch`

`rviz_robot.launch` is used to start a GUI that is used to locate the TurtleBot and to create NavGoals which tell the TurtleBot where to drive to. The GUI with buttons for leonardo is depicted in Figure 6.2. Buttons for additional robots and functions can be added with the plus symbol. the checkbox next to leonardo on the left side of the screenshot has to be clicked to show the model of the robot. These steps have

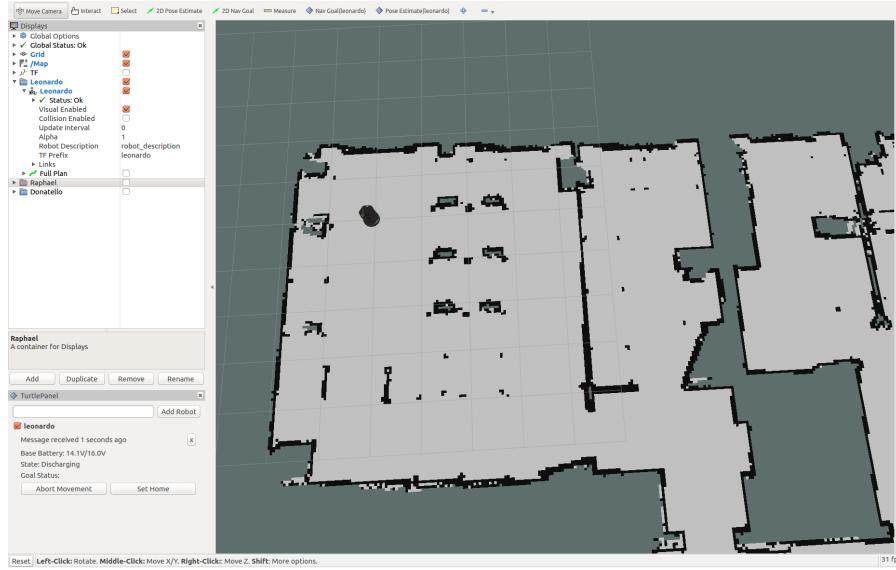


Figure 6.2: Rviz screen.

to be done before starting to processes on the control laptop. To start the TurtleBot four terminals with a ssh connection to the laptop on the TurtleBot are needed, which can be done either with `ssh tb@robotname.local` or `ssh tb@robotIP`. Once the ssh connections have been established the following commands have to be executed on the terminals:

- `roscore`
- `rosrun ttb_udp_proxy ttb_udp_proxy`
- `roslaunch turtlebot Bringup ninjaturtles.launch`
- `roslaunch turtlebot Bringup amcl.launch`

`ninjaturtles.launch` starts all parts of the TurtleBot, `amcl.launch` enables the motion and `ttb_udp_proxy` is used to communicate with other laptops/pcs. As soon

as all processes have been started the TurtleBot's location has to be set on rviz regarding the position in the real world and its orientation. This is done by clicking on the pose estimation button with the specific robot name. Afterwards the position and orientation of the TurtleBot can be determined by clicking and holding the left mouse button on the position the TurtleBot should be placed and dragging the cursor in the direction it should face. Once the left mouse button is released position and orientation are determined. To send the TurtleBot to a position the Nav Goal button is used. After clicking this button a position on the map can be marked as a navigation goal in the same way the position of the TurtleBot was determined.

6.2 Drive Topic and Message

To steer the TurtleBot without using amcl you need to publish a message under a specific topic. A message is published by using `rostopic pub TOPIC MESSAGE`. The topic used to steer the TurtleBot is named `/cmd_vel_mux/input/teleop` the message is named `geometry_msgs/Twist`. After the command is entered the message contents can be show by pressing tab twice and can be edited before sending the message with enter. Linear x is used to move the robot forward with a max value of 1.5. Angular z is used to rotate the robot. The max value is 14.

7 TurtleBot Gazebo Simulation

7.1 .bashrc

```
export TURTLEBOT_STACKS = ninja-hexagons
export TURTLEBOT_3D_SENSOR = ninja-kinect
export GAZEBO_MODEL_PATH
```

7.2 Starting Gazebo simulation

```
roslaunch turtlebot_bringup simulation.launch
Starts Gazebo, the map server and the robot
```

```
roslaunch turtlebot_bringup navigation_gazebo.launch
Starts amcl move_base
```

```
roslaunch rviz_sim.launch
Starts local cost map
```

7.3 topics

```
/leonardo/move_base_simple/goal
to publish the robots destination
```

```
/leonardo/initialpose pose estimate
```